

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: moons = pd.DataFrame( np.load("moon-all-input.npy") )
moons.columns = ["x1", "x2"]
moons
```

```
Out[ ]:
```

	x1	x2
0	1.538699	0.188744
1	0.394699	0.323724
2	-0.919147	1.311882
3	1.053964	-0.700408
4	1.040678	-0.437339
...
495	1.163768	-0.710319
496	0.477238	-0.681340
497	-0.246390	0.709156
498	0.897480	0.382935
499	0.413784	0.498548

500 rows × 2 columns

```
In [ ]: moons["category"] = pd.DataFrame(np.load("moon-all-output.npy"))
```

```
In [ ]: moons
```


```
Out[ ]:
```

	x1	x2	category
0	1.538699	0.188744	1
1	0.394699	0.323724	0
2	-0.919147	1.311882	0
3	1.053964	-0.700408	1
4	1.040678	-0.437339	1
...
495	1.163768	-0.710319	1
496	0.477238	-0.681340	1
497	-0.246390	0.709156	0
498	0.897480	0.382935	0
499	0.413784	0.498548	0

500 rows × 3 columns

```
In [ ]: sns.set_theme()
sns.color_palette("Set2")
sns.scatterplot(data = moons, x="x1", y="x2", hue = "category", style = "category")
```

```
Out[ ]: <Axes: xlabel='x1', ylabel='x2'>
```



```
In [ ]: moon_train, moon_test = moons.iloc[:375], moons.iloc[375:]  
moon_train
```

Out[]:

	x1	x2	category
0	1.538699	0.188744	1
1	0.394699	0.323724	0
2	-0.919147	1.311882	0
3	1.053964	-0.700408	1
4	1.040678	-0.437339	1
...
370	0.435785	-0.186108	1
371	-0.565615	0.563578	0
372	-1.205374	1.285517	0
373	0.515727	-0.347822	1
374	0.561705	0.388663	0

375 rows × 3 columns

In []: moon_test

```
Out[ ]:
```

	x1	x2	category
375	0.501693	0.217172	1
376	0.227405	0.589937	1
377	0.268167	0.431114	1
378	1.312708	-0.595236	1
379	0.418356	-0.382677	1
...
495	1.163768	-0.710319	1
496	0.477238	-0.681340	1
497	-0.246390	0.709156	0
498	0.897480	0.382935	0
499	0.413784	0.498548	0

125 rows × 3 columns

```
In [ ]: X_train = moon_train.iloc[:, :2].copy()
y_train = moon_train.iloc[:, 2].copy()

X_test = moon_test.iloc[:, :2].copy()
y_test = moon_test.iloc[:, 2].copy()
```

The data is largely balanced with the ratio being 0.344 instead of 0.333

```
In [ ]: moon_train["category"].sum(), moon_test["category"].sum()
```

```
Out[ ]: (186, 64)
```

AdaBoost

```
In [ ]: from sklearn.tree import DecisionTreeClassifier

import warnings

# Suppress FutureWarning messages
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```

class Classifier():
    classifier = None
    sample_weights = None
    error_weight = None
    classifier_weight = None
    learning_weight = 1
    X = None
    y = None
    y_pred = None

    def __init__(self, X, y, sample_weights, learning_rate=1):
        self.sample_weights = sample_weights
        self.learning_weight = learning_rate
        self.classifier = DecisionTreeClassifier(max_depth=1, random_state=42)
        self.X = X
        self.y = y

        self.classifier.fit(
            X, y, sample_weight=np.array(sample_weights).ravel())

    def calc_error(self):
        self.y_pred = self.classifier.predict(self.X)
        misclassified = self.sample_weights[self.y_pred != self.y]

        self.error_weight = misclassified.sum() / self.sample_weights.sum()

        return self.error_weight

    def calc_classifier_weight(self):
        self.calc_error()
        self.classifier_weight = (
            self.learning_weight * np.log((1-self.error_weight) / self.error_weight)).item()

        return self.classifier_weight

    def get_classifier_weight(self):
        return self.classifier_weight

    def get_new_sample_weights(self):
        self.calc_classifier_weight()

        current_weights = np.array(self.sample_weights)
        new_weights = np.empty(len(self.y_pred))
        pos_factor = np.exp(self.classifier_weight)
        neg_factor = np.exp(-self.classifier_weight)

        for i in range(len(self.y_pred)):

```

```

        if self.y_pred[i] != self.y[i]:
            new_weights[i] = current_weights[i] * pos_factor
        else:
            new_weights[i] = current_weights[i]
            # Slides contradict whether this should be the original value or the original * neg_factor
            # new_weights[i] = current_weights[i] * neg_factor

    # Normalizing the weights
    new_weights = new_weights / np.sum(new_weights)
    return new_weights

def predict(self, X):
    return self.classifier.predict(X).item()

```

In []: **class** AdaBoost():

```

    n_estimators = None
    classifiers = None
    X = None
    y = None
    learning_rate = 1

    def __init__(self, X, y, n_estimators=1, learning_rate=1):
        self.n_estimators = n_estimators
        self.classifiers = [None] * n_estimators
        self.X = X
        self.y = y
        self.learning_rate = learning_rate
        self.weights = None
        self.error_rates = None

    def fit(self):
        # Initializes the weights for all samples to 1/N, where N is the number of data points
        init_sample_weights = np.array(len(self.X) * [1/len(self.X)])
        self.classifiers[0] = Classifier(self.X, self.y, init_sample_weights)
        self.classifiers[0].calc_classifier_weight()

        for i in range(1, self.n_estimators):
            self.classifiers[i] = Classifier(
                self.X, self.y, self.classifiers[i-1].get_new_sample_weights())

        self.classifiers[-1].calc_classifier_weight()

        self.weights = np.array(
            [clf.classifier_weight for clf in self.classifiers])

        self.error_rates = np.array(

```

```

        [clf.error_weight for clf in self.classifiers])

def predict_sample(self, X, start, stop):
    voting_clfs = self.classifiers[start:stop]
    weight = self.weights[start:stop]
    # Calculate classifier weights for all classifiers at once

    # Predictions for all classifiers at once
    predictions = np.array([clf.predict(X) for clf in voting_clfs])

    # Count votes for class 1 and class 0
    vote_1 = np.sum(weight * (predictions == 1))
    vote_0 = np.sum(weight * (predictions == 0))

    if vote_1 > vote_0:
        return 1
    else:
        return 0

def predict(self, X, clf_start=0, clf_stop=3000):
    # Predicts with the full ensemble of n_estimators number of models
    stop = self.n_estimators

    if clf_stop < stop:
        stop = clf_stop

    predictions = np.empty(len(X), dtype=np.int8)

    for i in range(len(X)):
        predictions[i] = self.predict_sample(
            X.iloc[[i]], clf_start, stop)

    return predictions

def predict_ensembles(self, X):
    # X = X_input.to_numpy()

    predictions = np.empty((self.n_estimators, len(X)))

    voting_predictions = np.empty((len(X), self.n_estimators))

    for sample_row in range(len(X)):
        for clf in range(self.n_estimators):
            voting_predictions[sample_row][clf] = self.classifiers[clf].predict(X.iloc[[sample_row]])

    # print("Initial investment finished.")

```



```

    for k in range(self.n_estimators):
        for sample in range(len(X)):

            # Count votes for class 1 and class 0
            vote_1 = np.dot(self.weights[:,k], (voting_predictions[sample][:k] == 1))
            vote_0 = np.dot(self.weights[:,k], (voting_predictions[sample][:k] == 0))

            if vote_1 > vote_0:
                predictions[k][sample] = 1
            else:
                predictions[k][sample] = 0

        # print(f"{k} Estimator finished.")

    return predictions

```

```

In [ ]: estimators = 3000
        ada = AdaBoost(X_train, y_train, n_estimators= estimators)

        ada.fit()
        print("Done fitting")

```

Done fitting

Using AdaBoost implementation to get predictions on the Training and Testing datasets

```

In [ ]: y_pred_test = ada.predict_ensembles(X_test)

```

```

In [ ]: y_pred_train = ada.predict_ensembles(X_train)

```

```

In [ ]: y_true_test = np.array(y_test)
        y_true_train = np.array(y_train)

```

Getting Accuracy Scores

```

In [ ]: from sklearn.metrics import accuracy_score

        def get_accuracy(y_true, y_predictions, stop):
            scores = np.empty(stop)

            for i in range(stop):
                scores[i] = 1 - accuracy_score(y_true, y_predictions[i])
            score_df = pd.DataFrame(scores)
            score_df.to_csv("scores.csv")


```

```
return scores
```

```
test_scores = get_accuracy(y_true_test, y_pred_test, estimators)  
train_scores = get_accuracy(y_true_train, y_pred_train, estimators)
```

```
In [ ]: x = np.arange(3000)  
  
sns.lineplot(x = x, y = test_scores, label = "Test")  
sns.lineplot(x = x, y = train_scores, label = "Train")  
plt.xlabel("Estimator Number")  
plt.ylabel("Prediction Error")  
plt.title("Prediction Error after 3000 Iterations")
```

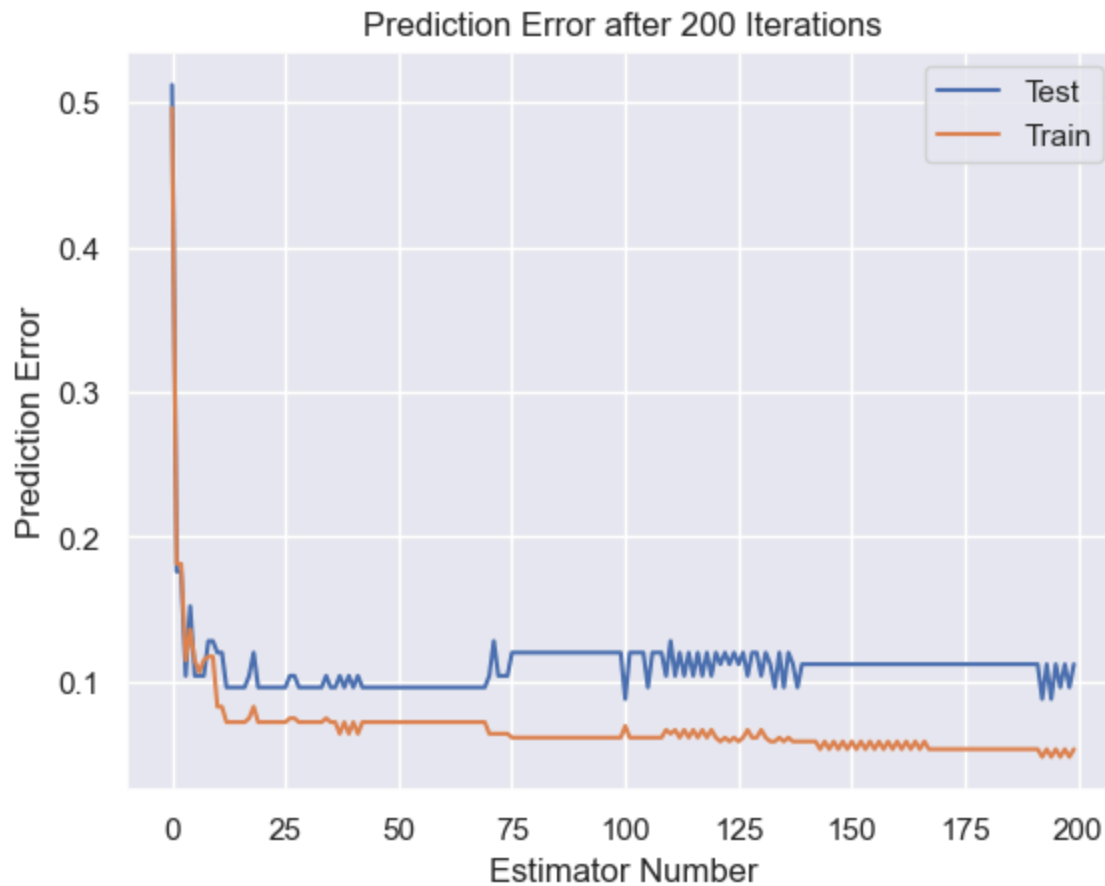
```
Out[ ]: Text(0.5, 1.0, 'Prediction Error after 3000 Iterations')
```



```
In [ ]: x = np.arange(200)  
  
sns.lineplot(x = x, y = test_scores[:200], label = "Test")  
sns.lineplot(x = x, y = train_scores[:200], label = "Train")
```

```
plt.xlabel("Estimator Number")
plt.ylabel("Prediction Error")
plt.title("Prediction Error after 200 Iterations")
```

Out[]: Text(0.5, 1.0, 'Prediction Error after 200 Iterations')

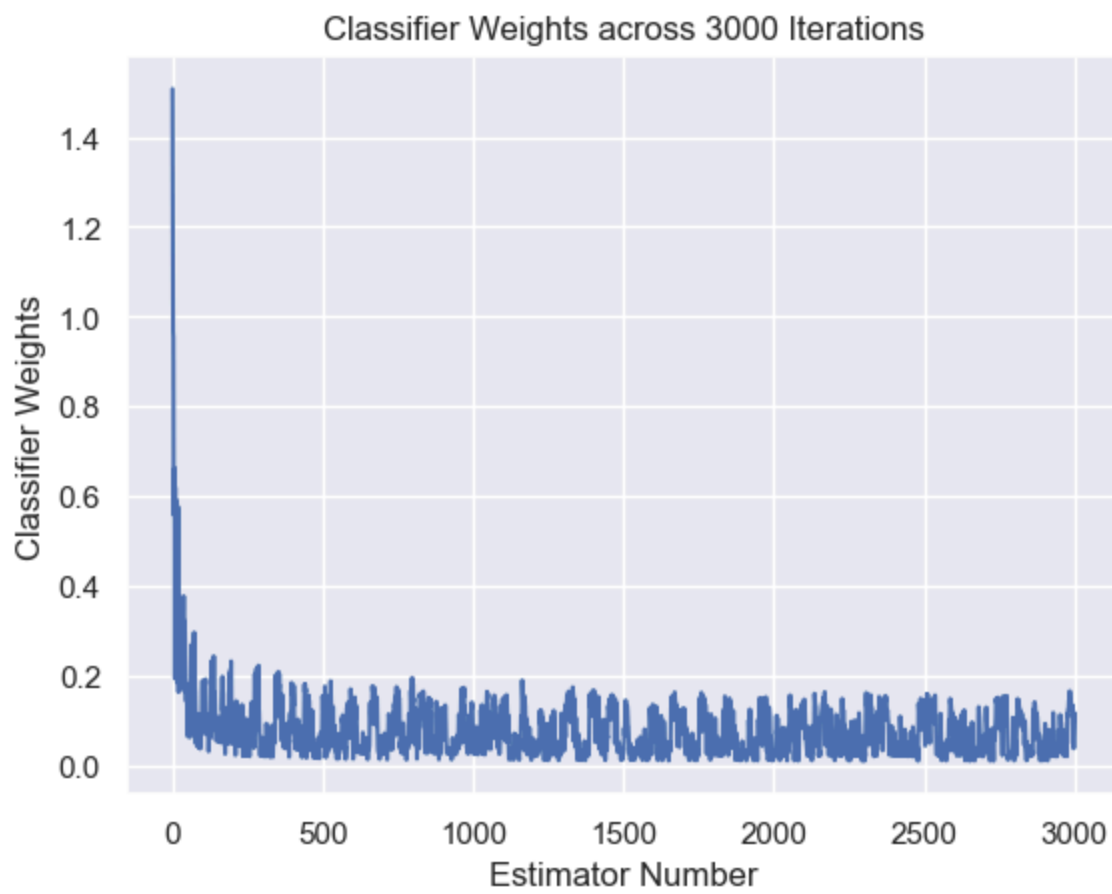


Observation: The model is able to reduce its training prediction error as the number of estimators increases; however this is at the cost of performance on the test data. This is a strong indicator of overfitting to the training dataset.

```
In [ ]: x = np.arange(3000)

sns.lineplot(x = x, y = ada.weights)
plt.xlabel("Estimator Number")
plt.ylabel("Classifier Weights")
plt.title("Classifier Weights across 3000 Iterations")
```

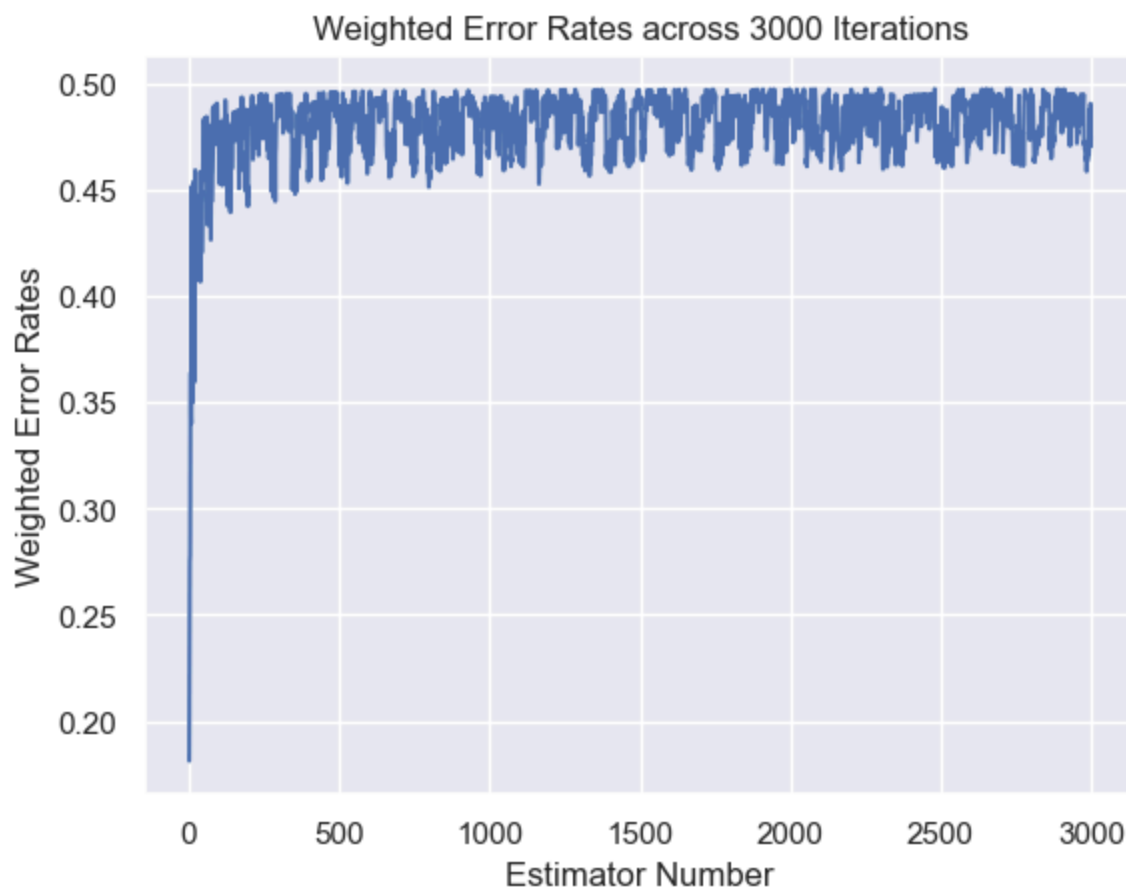
Out[]: Text(0.5, 1.0, 'Classifier Weights across 3000 Iterations')



```
In [ ]: x = np.arange(3000)

sns.lineplot(x = x, y = ada.error_rates)
plt.xlabel("Estimator Number")
plt.ylabel("Weighted Error Rates")
plt.title("Weighted Error Rates across 3000 Iterations")
```

```
Out[ ]: Text(0.5, 1.0, 'Weighted Error Rates across 3000 Iterations')
```



```
In [ ]: x = np.arange(3000)
scores = 1- train_scores

sns.lineplot(x = x, y = scores)
plt.xlabel("Estimator Number")
plt.ylabel("Accuracy")
plt.title("Accuracy on Training Set after 3000 Iterations")
```

```
Out[ ]: Text(0.5, 1.0, 'Accuracy on Training Set after 3000 Iterations')
```



```
In [ ]: x = np.arange(200)


sns.lineplot(x = x, y = scores[:200])
plt.xlabel("Estimator Number")
plt.ylabel("Accuracy")
plt.title("Accuracy on Training Set after 200 Iterations")
```

```
Out[ ]: Text(0.5, 1.0, 'Accuracy on Training Set after 200 Iterations')
```



Plotting Decision Boundaries

```
In [ ]: sns.scatterplot(data = moon_train, x = "x1", y = "x2", hue = "category", style = "category")
y_min, y_max = plt.gca().get_ylim()
x_min, x_max = plt.gca().get_xlim()
```



```
In [ ]: from sklearn.tree import plot_tree

indices = [0, 1, 2, 3, 4, 14, 20, 50, 100, 1000]

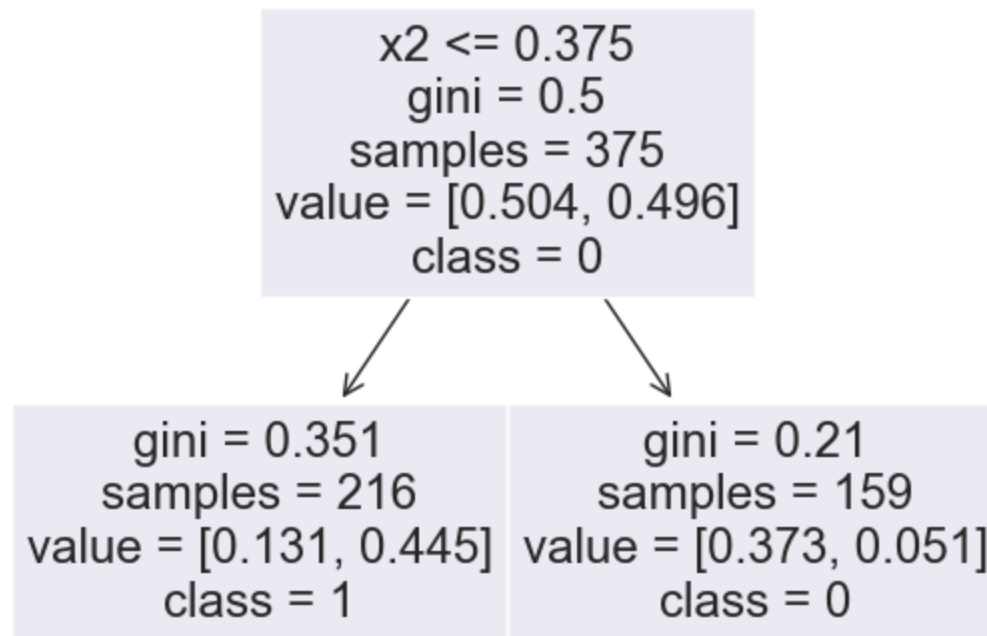
for index in indices:
    plt.clf()
    plot_tree(ada.classifiers[index].classifier, feature_names= ["x1", "x2"], class_names = ["0", "1"])
    plt.show()


    plt.clf()
    threshold = ada.classifiers[index].classifier.tree_.threshold[0]
    feature = ada.classifiers[index].classifier.tree_.feature[0]
    weight = ada.classifiers[index].sample_weights
    if feature == 1:
        plt.axhspan(ymin= y_min, ymax = threshold, color = "navajowhite", alpha = 0.5)
        plt.axhspan(ymin= threshold, ymax = y_max, color = "cornflowerblue", alpha = 0.3)
    else:
        plt.axvspan(xmin = x_min, xmax = threshold, color = "cornflowerblue", alpha = 0.3)
        plt.axvspan(xmin= threshold, xmax = x_max, color = "navajowhite", alpha = 0.5)

sns.scatterplot(data = moon_train, x = "x1", y= "x2", hue = "category", style = "category", size = weight)
```



```
plt.gca().set_xlim(x_min, x_max)  
plt.gca().set_ylim(y_min, y_max)  
plt.show()
```






$x_1 \leq 1.275$
gini = 0.484
samples = 375
value = [0.588, 0.412]
class = 0



gini = 0.408
samples = 295
value = [0.576, 0.231]
class = 0

gini = 0.119
samples = 80
value = [0.012, 0.181]
class = 1




x1 ≤ -0.551
gini = 0.482
samples = 375
value = [0.406, 0.594]
class = 1

```
graph TD; A["x1 ≤ -0.551  
gini = 0.482  
samples = 375  
value = [0.406, 0.594]  
class = 1"] --> B["gini = 0.0  
samples = 58  
value = [0.13, 0.0]  
class = 0"]; A --> C["gini = 0.433  
samples = 317  
value = [0.275, 0.594]  
class = 1"];
```

gini = 0.0
samples = 58
value = [0.13, 0.0]
class = 0

gini = 0.433
samples = 317
value = [0.275, 0.594]
class = 1




$x_2 \leq -0.343$
gini = 0.484
samples = 375
value = [0.59, 0.41]
class = 0



gini = 0.047
samples = 73
value = [0.003, 0.135]
class = 1

gini = 0.435
samples = 302
value = [0.587, 0.275]
class = 0




x1 ≤ -0.551
gini = 0.485
samples = 375
value = [0.412, 0.588]
class = 1



gini = 0.0
samples = 58
value = [0.066, 0.0]
class = 0

gini = 0.467
samples = 317
value = [0.346, 0.588]
class = 1




$x_1 \leq 1.578$
gini = 0.497
samples = 375
value = [0.536, 0.464]
class = 0

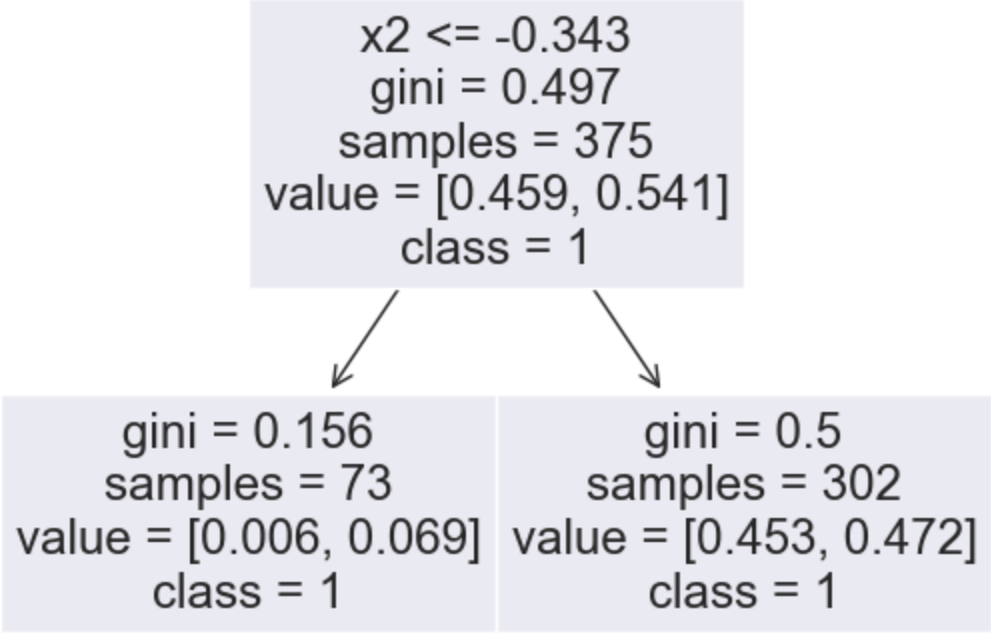


gini = 0.491
samples = 322
value = [0.536, 0.41]
class = 0

gini = 0.0
samples = 53
value = [0.0, 0.054]
class = 1




$x_2 \leq -0.343$
gini = 0.497
samples = 375
value = [0.459, 0.541]
class = 1



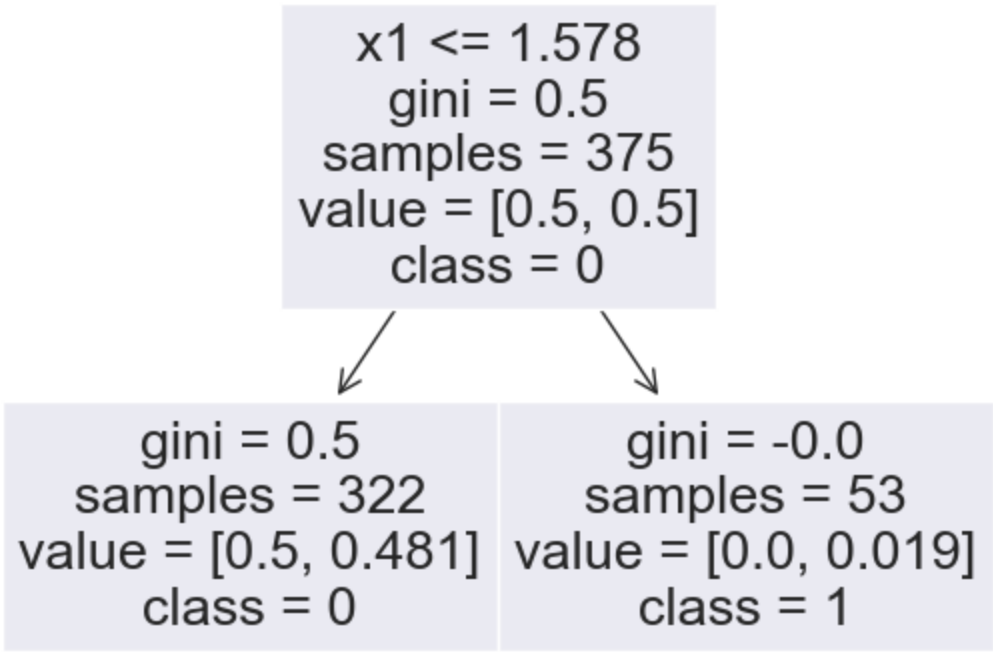
```
graph TD; A["x2 ≤ -0.343  
gini = 0.497  
samples = 375  
value = [0.459, 0.541]  
class = 1"] --> B["gini = 0.156  
samples = 73  
value = [0.006, 0.069]  
class = 1"]; A --> C["gini = 0.5  
samples = 302  
value = [0.453, 0.472]  
class = 1"];
```

gini = 0.156
samples = 73
value = [0.006, 0.069]
class = 1

gini = 0.5
samples = 302
value = [0.453, 0.472]
class = 1




$x_1 \leq 1.578$
gini = 0.5
samples = 375
value = [0.5, 0.5]
class = 0



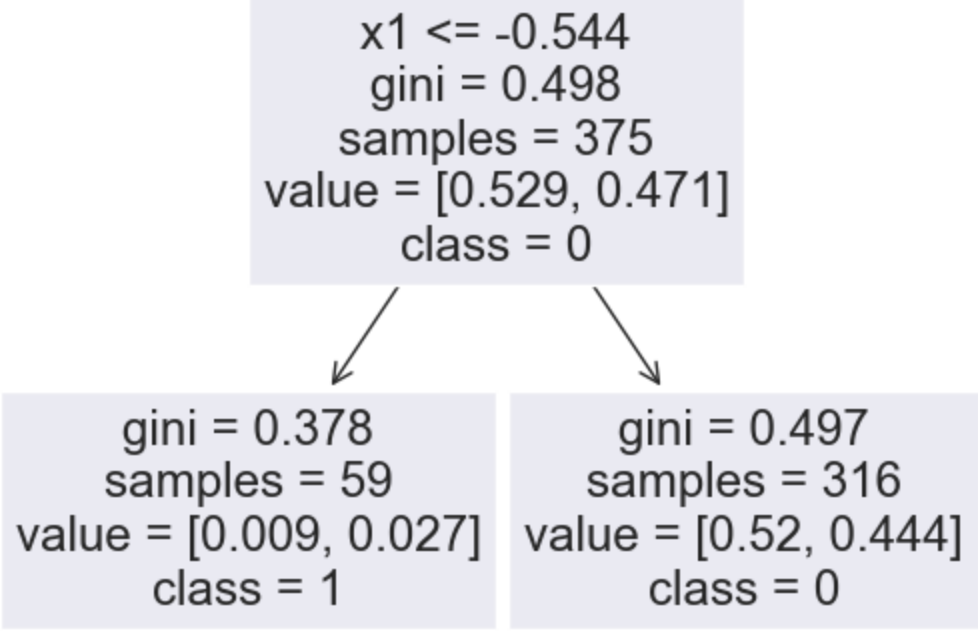
```
graph TD; A["x1 <= 1.578<br/>gini = 0.5<br/>samples = 375<br/>value = [0.5, 0.5]<br/>class = 0"] --> B["gini = 0.5<br/>samples = 322<br/>value = [0.5, 0.481]<br/>class = 0"]; A --> C["gini = -0.0<br/>samples = 53<br/>value = [0.0, 0.019]<br/>class = 1"];
```

gini = 0.5
samples = 322
value = [0.5, 0.481]
class = 0

gini = -0.0
samples = 53
value = [0.0, 0.019]
class = 1




x1 ≤ -0.544
gini = 0.498
samples = 375
value = [0.529, 0.471]
class = 0



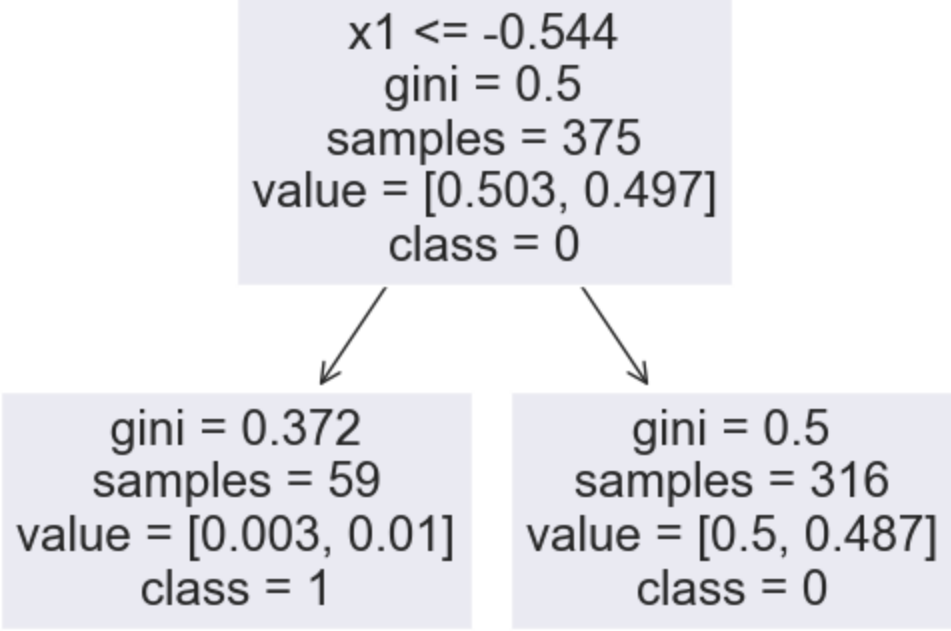
```
graph TD; A["x1 ≤ -0.544  
gini = 0.498  
samples = 375  
value = [0.529, 0.471]  
class = 0"] --> B["gini = 0.378  
samples = 59  
value = [0.009, 0.027]  
class = 1"]; A --> C["gini = 0.497  
samples = 316  
value = [0.52, 0.444]  
class = 0"];
```

gini = 0.378
samples = 59
value = [0.009, 0.027]
class = 1

gini = 0.497
samples = 316
value = [0.52, 0.444]
class = 0




x1 ≤ -0.544
gini = 0.5
samples = 375
value = [0.503, 0.497]
class = 0



```
graph TD; A["x1 ≤ -0.544  
gini = 0.5  
samples = 375  
value = [0.503, 0.497]  
class = 0"] --> B["gini = 0.372  
samples = 59  
value = [0.003, 0.01]  
class = 1"]; A --> C["gini = 0.5  
samples = 316  
value = [0.5, 0.487]  
class = 0"];
```

gini = 0.372
samples = 59
value = [0.003, 0.01]
class = 1

gini = 0.5
samples = 316
value = [0.5, 0.487]
class = 0



Plotting Decision Boundaries for Ensemble Models creating from AdaBoosting

```
In [ ]: # Basic meshgrid test

xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.3),
                     np.arange(y_min, y_max, 0.3))

# Combine xx and yy into a single DataFrame
df_grid = pd.DataFrame({'x1': xx.ravel(), 'x2': yy.ravel()})


indices = [0, 1, 2, 3, 4, 10, 14, 20, 50, 100, 200, 500, 1000]

for index in indices:
    # Predict the labels for all the points in the grid
    Z = ada.predict(df_grid, clf_stop = index)


    # Reshape Z to the shape of xx and yy
    Z = Z.reshape(xx.shape)

    # Plot the decision boundary and the data points
```


```
plt.title(f"Decision Boundary for Ensemble Model {index}")  
plt.contourf(xx, yy, Z, alpha=0.4)  
sns.scatterplot(data = moon_train, x = "x1", y = "x2", hue = "category", style = "category")  
plt.show()
```




Decision Boundary for Ensemble Model 1




Decision Boundary for Ensemble Model 2




Decision Boundary for Ensemble Model 3




Decision Boundary for Ensemble Model 4




Decision Boundary for Ensemble Model 10




Decision Boundary for Ensemble Model 14




Decision Boundary for Ensemble Model 20




Decision Boundary for Ensemble Model 50




Decision Boundary for Ensemble Model 100




Decision Boundary for Ensemble Model 200



Decision Boundary for Ensemble Model 500





Discussion

I have gained a deeper understanding of how ensemble models can be made more powerful from several weaker models. I can easily see how flexible this machine learning technique is through the use of more than just DecisionTree stumps to create a model that takes advantage of several machine learning algorithms at once. Moreover, the plotting of accuracy for the training and test dataset demonstrates the dangers of overfitting and letting a model generalize.